

# Untitled

October 14, 2022

```
[1]: # Load libs
import os
import PIL.Image

import numpy as np
import pandas as pd
import seaborn
import matplotlib.pyplot as plt

from tensorflow.keras.preprocessing.image import ImageDataGenerator, \
    img_to_array

from tensorflow.keras.layers import Dense, Dropout, GlobalAveragePooling2D, \
    Input, Rescaling

from tensorflow.keras.losses import SparseCategoricalCrossentropy
from tensorflow.keras.models import Model
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint
from tensorflow.keras.applications import MobileNetV3Small
```

2022-10-14 21:13:24.187679: I tensorflow/core/util/util.cc:169] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable `TF\_ENABLE\_ONEDNN\_OPTS=0`.

2022-10-14 21:13:24.244628: W tensorflow/stream\_executor/platform/default/dso\_loader.cc:64] Could not load dynamic library 'libcudart.so.11.0'; dlderror: libcudart.so.11.0: cannot open shared object file: No such file or directory

2022-10-14 21:13:24.244646: I tensorflow/stream\_executor/cuda/cudart\_stub.cc:29] Ignore above cudart dlerror if you do not have a GPU set up on your machine.

```
[3]: # Create a list of image pathes
images = []
for i, j, s in os.walk("."):
    print(i)
    for im in s:
        images.append(i+"/"+im)
```

```
.  
./ipynb_checkpoints  
./data  
./data/Sandal  
./data/Boot  
./data/Shoe  
./mobnet  
./mobnet/assets  
./mobnet/variables
```

```
[4]: # Check number of imgs  
len(images)
```

```
[4]: 15006
```

```
[5]: df = pd.DataFrame(  
    data = {"paths": images[2:], "ctg": [p.split("/")[2] for p in images[2:]]},  
    )  
df.head()
```

```
[5]:
```

|   | paths                           | ctg    |
|---|---------------------------------|--------|
| 0 | ./data/Sandal/Sandal (3877).jpg | Sandal |
| 1 | ./data/Sandal/Sandal (3032).jpg | Sandal |
| 2 | ./data/Sandal/Sandal (1653).jpg | Sandal |
| 3 | ./data/Sandal/Sandal (1008).jpg | Sandal |
| 4 | ./data/Sandal/Sandal (2139).jpg | Sandal |

```
[6]: # Shuffle it  
df = df.sample(frac = 1).reset_index()
```

```
[7]: # Show some examples of imgs  
fig = plt.figure(figsize = (14, 9))  
for i in range(20):  
    plt.subplot(5, 4, i+1)  
    img = PIL.Image.open(df.loc[i, "paths"])  
    plt.imshow(img)  
    plt.axis("off")  
    plt.title(df.loc[i, "ctg"])  
    plt.rcParams["figure.facecolor"] = "white"  
img.close()
```



```
[8]: batch_size = 32
```

## 1 Image data augmentation

```
[9]: data_gen = ImageDataGenerator(
    zoom_range = [0.8, 1.3],
    width_shift_range = 0.2,
    horizontal_flip = True,
    rotation_range = 20,
    validation_split = 0.25
)
```

```
[10]: # create a train data generator
train_data = data_gen.flow_from_directory(
    "data",
    target_size=(256, 256),
    batch_size = batch_size,
    shuffle = True,
    class_mode = "binary",
    subset = "training"
```

```
)
```

Found 11250 images belonging to 3 classes.

```
[11]: # create a validation data generator
val_data = data_gen.flow_from_directory(
    "data",
    target_size=(256, 256),
    batch_size = batch_size,
    shuffle = False,
    class_mode = "binary",
    subset = "validation"
)
```

Found 3750 images belonging to 3 classes.

```
[12]: def transfer_learn(premade):
    x = GlobalAveragePooling2D()(premade.output)
    x = Dense(1024, activation = "relu")(x)
    x = Dropout(0.2)(x)
    x = Dense(512, activation = "relu")(x)
    x = Dropout(0.2)(x)
    out = Dense(3, activation = "softmax")(x)
    return Model(inputs = premade.input, outputs = out)
```

```
[13]: # load pretrained mobilenet
mobilenet = MobileNetV3Small(
    include_top = False,
    input_shape = (256, 256, 3),
    weights = "imagenet"
)
mobilenet.trainable = False
```

WARNING:tensorflow:`input\_shape` is undefined or non-square, or `rows` is not 224. Weights for input shape (224, 224) will be loaded as the default.

```
2022-10-14 21:13:27.941337: E
tensorflow/stream_executor/cuda/cuda_driver.cc:271] failed call to cuInit:
CUDA_ERROR_COMPAT_NOT_SUPPORTED_ON_DEVICE: forward compatibility was attempted
on non supported HW
2022-10-14 21:13:27.941386: I
tensorflow/stream_executor/cuda/cuda_diagnostics.cc:169] retrieving CUDA
diagnostic information for host: sady
2022-10-14 21:13:27.941395: I
tensorflow/stream_executor/cuda/cuda_diagnostics.cc:176] hostname: sady
2022-10-14 21:13:27.941520: I
tensorflow/stream_executor/cuda/cuda_diagnostics.cc:200] libcuda reported
```

```

version is: 515.76.0
2022-10-14 21:13:27.941551: I
tensorflow/stream_executor/cuda/cuda_diagnostics.cc:204] kernel reported version
is: 515.65.1
2022-10-14 21:13:27.941560: E
tensorflow/stream_executor/cuda/cuda_diagnostics.cc:313] kernel version 515.65.1
does not match DSO version 515.76.0 -- cannot find working devices in this
configuration
2022-10-14 21:13:27.941952: I tensorflow/core/platform/cpu_feature_guard.cc:193]
This TensorFlow binary is optimized with oneAPI Deep Neural Network Library
(oneDNN) to use the following CPU instructions in performance-critical
operations: AVX2 AVX512F AVX512_VNNI FMA
To enable them in other operations, rebuild TensorFlow with the appropriate
compiler flags.

```

```

[14]: # get our final model
model1 = transfer_learn(mobilenet)

```

```

[15]: epochs = 100

```

```

[16]: model1.compile(
    optimizer = Adam(learning_rate = 0.001),
    loss = SparseCategoricalCrossentropy(from_logits = False),
    metrics=["accuracy"]
)

```

```

[17]: callbacks = [EarlyStopping(monitor = "val_accuracy", patience = 2, min_delta = 0.01),
    ModelCheckpoint("mobnet", monitor = "val_accuracy")]

```

```

[18]: # Train it
history = model1.fit(
    train_data,
    validation_data = val_data,
    batch_size = batch_size,
    epochs = epochs,
    callbacks = callbacks
)

```

Epoch 1/100

352/352 [=====] - ETA: 0s - loss: 0.2446 - accuracy: 0.9043

WARNING:absl:Found untraced functions such as \_jit\_compiled\_convolution\_op, \_jit\_compiled\_convolution\_op, \_jit\_compiled\_convolution\_op, \_jit\_compiled\_convolution\_op while saving (showing 5 of 52). These functions will not be directly callable after loading.

INFO:tensorflow:Assets written to: mobnet/assets

```

INFO:tensorflow:Assets written to: mobnet/assets

352/352 [=====] - 223s 625ms/step - loss: 0.2446 -
accuracy: 0.9043 - val_loss: 0.1361 - val_accuracy: 0.9507
Epoch 2/100
352/352 [=====] - ETA: 0s - loss: 0.1505 - accuracy:
0.9460

WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op,
_jit_compiled_convolution_op, _jit_compiled_convolution_op,
_jit_compiled_convolution_op, _jit_compiled_convolution_op while saving (showing
5 of 52). These functions will not be directly callable after loading.

INFO:tensorflow:Assets written to: mobnet/assets

INFO:tensorflow:Assets written to: mobnet/assets

352/352 [=====] - 226s 641ms/step - loss: 0.1505 -
accuracy: 0.9460 - val_loss: 0.1315 - val_accuracy: 0.9475
Epoch 3/100
352/352 [=====] - ETA: 0s - loss: 0.1298 - accuracy:
0.9534

WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op,
_jit_compiled_convolution_op, _jit_compiled_convolution_op,
_jit_compiled_convolution_op, _jit_compiled_convolution_op while saving (showing
5 of 52). These functions will not be directly callable after loading.

INFO:tensorflow:Assets written to: mobnet/assets

INFO:tensorflow:Assets written to: mobnet/assets

352/352 [=====] - 224s 636ms/step - loss: 0.1298 -
accuracy: 0.9534 - val_loss: 0.1800 - val_accuracy: 0.9395

```

## 1.1 Evaluate model

```
[19]: from sklearn.metrics import confusion_matrix, classification_report
```

```
[20]: y_pred = model1.predict(val_data, 32)
      y_pred = np.argmax(y_pred, axis = 1)
      y_pred
```

```
118/118 [=====] - 54s 450ms/step
```

```
[20]: array([0, 0, 0, ..., 2, 2, 2])
```

```
[21]: y_true = val_data.classes
      y_true
```

```
[21]: array([0, 0, 0, ..., 2, 2, 2], dtype=int32)
```


```
[22]: lbls = val_data.class_indices.keys()
```

```
[23]: cm = confusion_matrix(y_true, y_pred, normalize = "true")
```

```
[24]: cm
```

```
[24]: array([[0.8768, 0.0456, 0.0776],  
          [0.0024, 0.972 , 0.0256],  
          [0.0016, 0.04  , 0.9584]])
```

```
[25]: lbls = val_data.class_indices.keys()  
      seaborn.heatmap(cm, xticklabels = lbls, yticklabels=lbls)  
      plt.show()
```



```
[26]: val_data.class_indices
```

```
[26]: {'Boot': 0, 'Sandal': 1, 'Shoe': 2}
```

```
[27]: print(  
      classification_report(y_true, y_pred, target_names = val_data.class_indices.  
      ↪keys())  
      )
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| Boot         | 1.00      | 0.88   | 0.93     | 1250    |
| Sandal       | 0.92      | 0.97   | 0.94     | 1250    |
| Shoe         | 0.90      | 0.96   | 0.93     | 1250    |
| accuracy     |           |        | 0.94     | 3750    |
| macro avg    | 0.94      | 0.94   | 0.94     | 3750    |
| weighted avg | 0.94      | 0.94   | 0.94     | 3750    |

```
[ ]:
```